Lab 8. Advanced Ansible Topics

ERNESTO .NET

In this lab, we will cover the following topics:

- Asynchronous versus synchronous actions
- Controlling play execution for rolling updates
- Configuring the maximum failure percentage
- Setting task execution delegation
- Using the [run_once] option
- Running playbooks locally
- Working with proxies and jump hosts
- Configuring playbook prompts
- Placing tags in the plays and tasks
- Securing data with Ansible Vault

Asynchronous versus synchronous actions

Ansible tasks can be run asynchronously---that is to say, tasks can be run in the background on the target host and polled on a regular basis. This is in contrast to synchronous tasks, where the connection to the target host is kept open until the task completes (which runs the risk of a timeout occurring).

As ever, let's explore this through a practical example. Suppose we have two servers in a simple INI-formatted inventory:

```
[frontends]
frt01.example.com
frt02.example.com
```

Now, in order to simulate a long-running task, we'll run the [sleep] command using the [shell] module. However, rather than have it run with the SSH connection blocked for the duration of the [sleep] command, we'll add two special parameters to the task, as shown:

```
---
- name: Play to demonstrate asynchronous tasks
  hosts: frontends
  become: true

  tasks:
    - name: A simulated long running task
      shell: "sleep 20"
      async: 30
      poll: 5
```

The two new parameters are [async] and [poll]. The [async] parameter tells Ansible that this task should be run asynchronously (so that the SSH connection will not be blocked) for a maximum of [30] seconds. If the task runs for longer than this configured time, Ansible considers the task to have failed and the play is failed, accordingly. When [poll] is set to a positive integer, Ansible checks the status of the asynchronous task at the specified interval---in this example, every [5] seconds. If [poll] is set to [0], then the task is run in the background and never checked---it is up to you to write a task to manually check its status later on.

If you don't specify the [poll] value, it will be set to the default value defined by the [DEFAULT_POLL_INTERVAL] configuration parameter of Ansible (which is [10] seconds).

When you run this playbook, you will find that it runs just like any other playbook; from the terminal output, you won't be able to see any difference. But behind the scenes, Ansible checks the task every [5] seconds until it succeeds or reaches the [async] timeout value of [30] seconds:

```
$ ansible-playbook -i hosts async.yml

PLAY [Play to demonstrate asynchronous tasks] *********************************

TASK [Gathering Facts] *******************************************************
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [A simulated long running task] *****************************************
changed: [frt02.example.com]
changed: [frt01.example.com]

PLAY RECAP *******************************************************************
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

If you want to check on the task later (that is, if [poll] is set to [0]), you could add a second task to your playbook so that it looks as follows:

```
---
- name: Play to demonstrate asynchronous tasks
  hosts: frontends
  become: true

  tasks:
    - name: A simulated long running task
      shell: "sleep 20"
      async: 30
      poll: 0
      register: long_task

    - name: Check on the asynchronous task
      async_status:
        jid: "{{ long_task.ansible_job_id }}"
      register: async_result
      until: async_result.finished
      retries: 30
```

In this playbook, the initial asynchronous task is defined as before, except we have now set [poll] to [0]. We have also chosen to register the result of this task to a variable called [long_task]---this is so that we can query the job ID for the task when we check it later on. The next (new) task in the play uses the [async_status] module to check on the job ID we registered from the first task and loops until the job either finishes or reaches [30] retries---whichever comes first. When using these in a playbook, you almost certainly wouldn't add the two tasks back to back like this--- usually, you would perform additional tasks in between them---but to keep this example simple, we will run the two tasks sequentially. Running this playbook should yield an output similar to the following:

```
$ ansible-playbook -i hosts async2.yml

PLAY [Play to demonstrate asynchronous tasks] *********************************

TASK [Gathering Facts] *******************************************************
ok: [frt01.example.com]
ok: [frt02.example.com]

TASK [A simulated long running task] *****************************************
changed: [frt02.example.com]
changed: [frt01.example.com]

TASK [Check on the asynchronous task] ****************************************
FAILED - RETRYING: Check on the asynchronous task (30 retries left).
FAILED - RETRYING: Check on the asynchronous task (30 retries left).
FAILED - RETRYING: Check on the asynchronous task (29 retries left).
FAILED - RETRYING: Check on the asynchronous task (29 retries left).
FAILED - RETRYING: Check on the asynchronous task (28 retries left).
FAILED - RETRYING: Check on the asynchronous task (28 retries left).
FAILED - RETRYING: Check on the asynchronous task (27 retries left).
FAILED - RETRYING: Check on the asynchronous task (27 retries left).
changed: [frt01.example.com]
changed: [frt02.example.com]

PLAY RECAP *******************************************************************
frt01.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

In the preceding code block, we can see that the long-running task is left running and the next task polls its status until the conditions we set are met. In this case, we can see that the task finished successfully and the overall play result was successful. Asynchronous actions are especially useful for large downloads, package updates, and other tasks that might take a long time to run. You may find them useful in your playbook development, especially in more complex infrastructures.

With this under our belt, let's take a look at another advanced technique that might be useful in large infrastructures---performing rolling updates with Ansible.

Control play execution for rolling updates

By default, Ansible parallelizes tasks on multiple hosts at the same time to speed up automation tasks in large inventories. The setting for this is defined by the [forks] parameter in the Ansible configuration file, which defaults to [5] (so, by default, Ansible attempts to run its automation job on five hosts at the same time).

1. Create the following simple playbook to run two commands on the two hosts in our inventory. The content of the command is not important at this stage, but if you run the [date] command using the [command] module, you will be able to see the time that each task is run, as well as if you specify [-v] to increase the verbosity when you run the play:

```
---
- name: Simple serial demonstration play
```

```
  hosts: frontends
  gather_facts: false

  tasks:
    - name: First task
      command: date
    - name: Second task
      command: date
```

2. Now, if you run this play, you will see that it performs all the operations on each host simultaneously, as we have fewer hosts than the default number of forks---[5]. This behavior is normal for Ansible, but not really what we want as our users will experience service outage:

```
$ ansible-playbook -i hosts serial.yml

PLAY [Simple serial demonstration play] ***************************************

TASK [First task] ************************************************************
changed: [frt02.example.com]
changed: [frt01.example.com]

TASK [Second task] ***********************************************************
changed: [frt01.example.com]
changed: [frt02.example.com]


PLAY RECAP *******************************************************************
frt01.example.com : ok=2 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

3. Now, let's modify the play definition, as shown. We'll leave the [tasks] sections exactly as they were in *step 1*:

```
---
- name: Simple serial demonstration play
  hosts: frontends
  serial: 1
  gather_facts: false
```

4. Notice the presence of the [serial: 1] line. This tells Ansible to complete the play on [1] host at a time before moving on to the next. If we run the play again, we can see this in action:

```
$ ansible-playbook -i hosts serial.yml

PLAY [Simple serial demonstration play] ***************************************

TASK [First task] ************************************************************
changed: [frt01.example.com]

TASK [Second task] ***********************************************************
changed: [frt01.example.com]
```

```
PLAY [Simple serial demonstration play] *****************************************

TASK [First task] **************************************************************
changed: [frt02.example.com]

TASK [Second task] *************************************************************
changed: [frt02.example.com]

PLAY RECAP *********************************************************************
frt01.example.com : ok=2 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Much better! If you imagine that this playbook actually disables these hosts on a load balancer, performs an upgrade, and then re-enables the hosts on the load balancer, this is exactly how you would want the operation to proceed. Doing so without the [serial: 1] directive would result in all the hosts being removed from the load balancer at once, causing a loss of service.

It is useful to note that the [serial] directive can also take a percentage instead of an integer. When you specify a percentage, you are telling Ansible to run the play on that percentage of hosts at one time. So, if you have [4] hosts in your inventory and specify [serial: 25%], Ansible will only run the play on one host at a time. If you have [8] hosts in your inventory, it will run the play on two hosts at a time. I'm sure you get the idea!

You can even build on this by passing a list to the [serial] directive. Consider the following code:

```
serial:
  - 1
  - 3
  - 5
```

This tells Ansible to run the play on [1] host, initially, then on the next [3], and then on batches of [5] at a time until the inventory is completed. You can also specify a list of percentages in place of the integer numbers of hosts. In doing this, you will build up a robust playbook that can perform rolling updates without causing a loss of service to end users. With this complete, let's further build on this knowledge by looking at controlling the maximum failure percentage that Ansible can tolerate before it aborts a play, which will again be useful in highly available or load-balanced environments such as this.

Configuring the maximum failure percentage

In its default mode of operation, Ansible continues to execute a play on a batch of servers (the batch size is determined by the [serial] directive we discussed in the preceding section) as long as there are hosts in the inventory and a failure isn't recorded. Obviously, in a highly available or load-balanced environment (such as the one we discussed previously), this is not ideal. If there is a bug in your play, or perhaps a problem with the code being rolled out, the last thing that you want is for Ansible to faithfully roll it out to all servers in the cluster, causing a service outage because all the nodes suffered a failed upgrade. It would be far better, in this kind of environment, to fail early on and leave at least some hosts in the cluster untouched until someone can intervene and resolve the issue.

For our practical example, let's consider an expanded inventory with [10] hosts in it. We'll define this as follows:

```
[frontends]
frt[01:10].example.com
```

Now, let's create a simple playbook to run on these hosts. We will set our batch size to [5] and [max_fail_percentage] to [50%] in the play definition:

1. Create the following play definition to demonstrate the use of the [max_fail_percentage] directive:

```
---
- name: A simple play to demonstrate use of max_fail_percentage
  hosts: frontends
  gather_facts: no
  serial: 5
  max_fail_percentage: 50
```

We have defined [10] hosts in our inventory, so it will process them in batches of 5 (as specified by [serial: 5]). We will fail the entire play and stop performing processing if more than 50% of the hosts in one batch fails.

The number of failed hosts must exceed the value of [max_fail_percentage]; if it is equal, the play continues. So, in our example, if exactly 50% of our hosts failed, the play would still continue.

2. Next, we will define two simple tasks. The first task has a special clause under it that we use to deliberately simulate a failure---this line starts with [failed_when] and we use it to tell the task that if it runs this task on the first three hosts in the batch, then it should deliberately fail this task regardless of the result; otherwise, it should allow the task to run as normal:

```
tasks:
  - name: A task that will sometimes fail
    debug:
      msg: This might fail
    failed_when: inventory_hostname in ansible_play_batch[0:3]
```

3. Finally, we'll add a second task that will always succeed. This is run if the play is allowed to continue, but not if it is aborted:

```
  - name: A task that will succeed
    debug:
      msg: Success!
```

So, we have deliberately constructed a playbook that will run on a 10-host inventory in batches of 5 hosts at a time, but the play is aborted if more than 50% of the hosts in any given batch experiences a failure. We have also deliberately set up a failure condition that causes three of the hosts in the first batch of 5 (60%) to fail.

4. Run the playbook and let's observe what happens:

```
$ ansible-playbook -i morehosts maxfail.yml

PLAY [A simple play to demonstrate use of max_fail_percentage] *****************

TASK [A task that will sometimes fail] ****************************************
fatal: [frt01.example.com]: FAILED! => {
    "msg": "This might fail"
}
fatal: [frt02.example.com]: FAILED! => {
    "msg": "This might fail"
}
```

```
fatal: [frt03.example.com]: FAILED! => {
    "msg": "This might fail"
}
ok: [frt04.example.com] => {
    "msg": "This might fail"
}
ok: [frt05.example.com] => {
    "msg": "This might fail"
}


NO MORE HOSTS LEFT *********************************************************

NO MORE HOSTS LEFT *********************************************************

PLAY RECAP *****************************************************************
frt01.example.com : ok=0 changed=0 unreachable=0 failed=1 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=0 changed=0 unreachable=0 failed=1 skipped=0 rescued=0
ignored=0
frt03.example.com : ok=0 changed=0 unreachable=0 failed=1 skipped=0 rescued=0
ignored=0
frt04.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt05.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Notice the results of this playbook. We deliberately failed three of the first batch of 5, exceeding the threshold for [max_fail_percentage] that we set. This immediately causes the play to abort and the second task is not performed on the first batch of 5. You will also notice that the second batch of 5, out of the 10 hosts, is never processed, so our play was truly aborted. This is exactly the behavior you would want to see to prevent a failed update from rolling out across a cluster. Through the careful use of batches and [max_fail_percentage], you can safely run automated tasks across an entire cluster without the fear of breaking the entire cluster in the event of an issue. In the next section, we will take a look at another feature of Ansible that can be incredibly useful when it comes to working with clusters---task delegation.

Setting task execution delegation

In every play we have run so far, we have assumed that all the tasks are executed on each host in the inventory in turn. However, what if you need to run one or two tasks on a different host? For example, we have talked about the concept of automating upgrades on clusters. Logically, however, we would want to automate the entire process, including the removal of each host in turn from the load balancer and its return after the task is completed.

Although we still want to run our play across our entire inventory, we certainly don't want to run the load balancer commands from those hosts. Let's once again explain this in more detail with a practical example. We'll reuse the two simple host inventories that we used earlier in this lab:

```
[frontends]
frt01.example.com
frt02.example.com
```

Now, to work on this, let's create two simple shell scripts in the same directory as our playbook. These are only examples as setting up a load balancer is beyond the scope of this course. However, imagine that you have a shell

script (or other executables) that you can call that can add and remove hosts to and from a load balancer:

1. For our example, let's create a script called [remove_from_loadbalancer.sh], which will contain the following:

```
#!/bin/sh
echo Removing $1 from load balancer...
```

2. We will also create a script called [add_to_loadbalancer.sh], which will contain the following:

```
#!/bin/sh
echo Adding $1 to load balancer...
```

Obviously, in a real-world example, there would be much more code in these scripts!

3. Now, let's create a playbook that will perform the logic we outlined here. We'll first create a very simple play definition (you are free to experiment with the [serial] and [max_fail_percentage] directives as you wish) and an initial task:

```
---
- name: Play to demonstrate task delegation
  hosts: frontends

  tasks:
    - name: Remove host from the load balancer
      command: ./remove_from_loadbalancer.sh {{ inventory_hostname }}
      args:
        chdir: "{{ playbook_dir }}"
      delegate_to: localhost
```

Notice the task structure---most of it will be familiar to you. We are using the [command] module to call the script we created earlier, passing the hostname from the inventory being removed from the load balancer to the script. We use the [chdir] argument with the [playbook_dir] magic variable to tell Ansible that the script is to be run from the same directory as the playbook.

The special part of this task is the [delegate_to] directive, which tells Ansible that even though we're iterating through an inventory that doesn't contain [localhost], we should run this action on [localhost] (we aren't copying the script to our remote hosts, so it won't run if we attempt to run it from there).

4. After this, we add a task where the upgrade work is carried out. This task has no [delegate_to] directive, and so it is actually run on the remote host from the inventory (as desired):

```
    - name: Deploy code to host
      debug:
        msg: Deployment code would go here....
```

5. Finally, we add the host back to the load balancer using the second script we created earlier. This task is almost identical to the first:

```
    - name: Add host back to the load balancer
      command: ./add_to_loadbalancer.sh {{ inventory_hostname }}
      args:
        chdir: "{{ playbook_dir }}"
      delegate_to: localhost
```

6. Let's see this playbook in action:

```
$ ansible-playbook -i hosts delegate.yml

PLAY [Play to demonstrate task delegation] *************************************

TASK [Gathering Facts] ********************************************************
ok: [frt01.example.com]
ok: [frt02.example.com]

TASK [Remove host from the load balancer] *************************************
changed: [frt02.example.com -> localhost]
changed: [frt01.example.com -> localhost]

TASK [Deploy code to host] ****************************************************
ok: [frt01.example.com] => {
    "msg": "Deployment code would go here...."
}
ok: [frt02.example.com] => {
    "msg": "Deployment code would go here...."
}

TASK [Add host back to the load balancer] *************************************
changed: [frt01.example.com -> localhost]
changed: [frt02.example.com -> localhost]

PLAY RECAP ********************************************************************
frt01.example.com : ok=4 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=4 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Notice how even though Ansible is working through the inventory (which doesn't feature [localhost]), the load balancer-related scripts are actually run from [localhost], while the upgrade task is performed directly on the remote host. This, of course, isn't the only thing you can do with task delegation, but it's a common example of a way that it can help you.

In truth, you can delegate any task to [localhost], or even another non-inventory host. You could, for example, run an [rsync] command delegated to [localhost] to copy files to remote hosts using a similar task definition to the previous one. This is useful because although Ansible has a [copy] module, it can't perform the advanced recursive [copy] and [update] functions that [rsync] is capable of.

Also, note that you can choose to use a form of shorthand notation in your playbooks (and roles) for [delegate_to], called [local_action]. This allows you to specify a task on a single line that would ordinarily be run with [delegate_to: localhost] added below it. Wrapping this all up into a second example, our playbook will look as follows:

```
---
- name: Second task delegation example
  hosts: frontends

  tasks:
  - name: Perform an rsync from localhost to inventory hosts
    local_action: command rsync -a /tmp/ {{ inventory_hostname }}:/tmp/target/
```

The preceding shorthand notation is equivalent to the following:

```
tasks:
  - name: Perform an rsync from localhost to inventory hosts
    command: rsync -a /tmp/ {{ inventory_hostname }}:/tmp/target/
    delegate_to: localhost
```

If we run this playbook, we can see that [local_action] does indeed run [rsync] from [localhost], enabling us to efficiently copy whole directory trees across to remote servers in the inventory:

```
$ ansible-playbook -i hosts delegate2.yml

PLAY [Second task delegation example] ******************************************

TASK [Gathering Facts] *********************************************************
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [Perform an rsync from localhost to inventory hosts] **********************
changed: [frt02.example.com -> localhost]
changed: [frt01.example.com -> localhost]

PLAY RECAP *********************************************************************
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

This concludes our look at task delegation, although as stated, these are just two common examples. I'm sure you can think up some more advanced use cases for this capability. Let's continue looking at controlling the flow of Ansible code by proceeding, in the next section, to look at the special [run_once] option.

Using the run_once option

When working with clusters, you will sometimes encounter a task that should only be executed once for the entire cluster. For example, you might want to upgrade the schema of a clustered database or issue a command to reconfigure a Pacemaker cluster that would normally be issued on one node and replicated to all other nodes by Pacemaker. You could, of course, address this with a special inventory with only one host in it, or even by writing a special play that references one host from the inventory, but this is inefficient and starts to make your code fragmented.

Instead, you can write your code as you normally would, but make use of the special [run_once] directive for any tasks you want to run only once on your inventory. For example, let's reuse the 10-host inventory that we defined earlier in this lab. Now, let's proceed to demonstrate this option, as follows:

1. Create the simple playbook as in the following code block. We're using a debug statement to display some output, but in real life, you would insert your script or command that performs your one-off cluster function here (for example, upgrading a database schema):

```
---
- name: Play to demonstrate the run_once directive
  hosts: frontends
```

```
  tasks:
    - name: Upgrade database schema
      debug:
        msg: Upgrading database schema...
      run_once: true
```

2. Now, let's run this playbook and see what happens:

```
$ ansible-playbook -i morehosts runonce.yml

PLAY [Play to demonstrate the run_once directive] *****************************

TASK [Gathering Facts] *******************************************************
ok: [frt02.example.com]
ok: [frt05.example.com]
ok: [frt03.example.com]
ok: [frt01.example.com]
ok: [frt04.example.com]
ok: [frt06.example.com]
ok: [frt08.example.com]
ok: [frt09.example.com]
ok: [frt07.example.com]
ok: [frt10.example.com]

TASK [Upgrade database schema] ***********************************************
ok: [frt01.example.com] => {
    "msg": "Upgrading database schema..."
}
---

PLAY RECAP *******************************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt03.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt04.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt05.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt06.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt07.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt08.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt09.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt10.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Notice that, just as desired, although the playbook was run on all 10 hosts (and, indeed, gathered facts from all 10 hosts), we only ran the upgrade task on one host.

3. It's important to note that the [run_once] option applies per batch of servers, so if we add [serial: 5] to our play definition (running our play in two batches of 5 on our inventory of 10 servers), the schema upgrade task actually runs twice! It runs once as requested, but once per batch of servers, not once for the entire inventory. Be careful of this nuance when working with this directive in a clustered environment.

Add [serial: 5] to your play definition and rerun the playbook. The output should appear as follows:

```
$ ansible-playbook -i morehosts runonce.yml

PLAY [Play to demonstrate the run_once directive] *****************************

TASK [Gathering Facts] ********************************************************
ok: [frt04.example.com]
ok: [frt01.example.com]
ok: [frt02.example.com]
ok: [frt03.example.com]
ok: [frt05.example.com]

TASK [Upgrade database schema] ************************************************
ok: [frt01.example.com] => {
    "msg": "Upgrading database schema..."
}

PLAY [Play to demonstrate the run_once directive] *****************************

TASK [Gathering Facts] ********************************************************
ok: [frt08.example.com]
ok: [frt06.example.com]
ok: [frt07.example.com]
ok: [frt10.example.com]
ok: [frt09.example.com]

TASK [Upgrade database schema] ************************************************
ok: [frt06.example.com] => {
    "msg": "Upgrading database schema..."
}

PLAY RECAP ********************************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt03.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt04.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt05.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt06.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

```
frt07.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt08.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt09.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt10.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

This is how the [run_once] option is designed to work---you can observe, in the preceding output, that our schema upgrade ran twice, which is probably not something we wanted! However, with this awareness, you should be able to take advantage of this option to control your playbook flow across clusters and still achieve the results you want. Let's now move away from cluster-related Ansible tasks and look at the subtle but important difference between running playbooks locally and running them on [localhost].

Running playbooks locally

It is important to note that when we talk about running a playbook locally with Ansible, it is not the same as talking about running it on [localhost]. If we run a playbook on [localhost], Ansible actually sets up an SSH connection to [localhost] (it doesn't differentiate its behavior or attempt to detect whether a host in the inventory is local or remote---it simply tries faithfully to connect).

Indeed, we can try creating a [local] inventory file with the following contents:

```
[local]
localhost
```

Now, if we attempt to run the [ping] module in an ad hoc command against this inventory, we see the following:

```
$ ansible -i localhosts -m ping all --ask-pass
The authenticity of host 'localhost (::1)' can't be established.
ECDSA key fingerprint is SHA256:DUwVxH+45432pSr9qsN8Av4l0KJJ+r5jTo123n3XGvZs.
ECDSA key fingerprint is MD5:78:d1:dc:23:cc:28:51:42:eb:fb:58:49:ab:92:b6:96.
Are you sure you want to continue connecting (yes/no)? yes
SSH password:
localhost | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
```

As you can see, Ansible set up an SSH connection that needed the host key to validate, as well as our SSH password. Now, although you could add the host key (as we did in the preceding code block), add key-based SSH authentication to your [localhost], and so on, there is a more direct way of doing this.

We can now modify our inventory so that it looks as follows:

```
[local]
localhost ansible_connection=local
```

We've added a special variable to our [localhost] entry---the [ansible_connection] variable---which defines which protocol is used to connect to this inventory host. So, we have told it to use a direct local connection instead of an SSH-based connectivity (which is the default).

It should be noted that this special value for the [ansible_connection] variable actually overrides the hostname you have put in your inventory. So, if we change our inventory to look as follows, Ansible will not even attempt to connect to the remote host called [frt01.example.com---]it will connect locally to the machine running the playbook (without SSH):

```
[local]
frt01.example.com ansible_connection=local
```

We can demonstrate this very simply. Let's first check for the absence of a test file in our local [/tmp] directory:

```
ls -l /tmp/foo
ls: cannot access /tmp/foo: No such file or directory
```

Now, let's run an ad hoc command to touch this file on all hosts in the new inventory we just defined:

```
$ ansible -i localhosts2 -m file -a "path=/tmp/foo state=touch" all
frt01.example.com | CHANGED => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": true,
    "dest": "/tmp/foo",
    "gid": 0,
    "group": "root",
    "mode": "0644",
    "owner": "root",
    "size": 0,
    "state": "file",
    "uid": 0
}
```

The command ran successfully, so let's see whether the test file is present on the local machine:

```
$ ls -l /tmp/foo
-rw-r--r-- 1 root root 0 Apr 24 16:28 /tmp/foo
```

It is! So, the ad hoc command did not attempt to connect to [frt01.example.com], even though this host name was in the inventory. The presence of [ansible_connection=local] meant that this command was run on the local machine without using SSH.

This ability to run commands locally without the need to set up SSH connectivity, SSH keys, and so on can be incredibly valuable, especially if you need to get things up and running quickly on your local machine. With this complete, let's take a look at how you can work with proxies and jump hosts using Ansible.

Working with proxies and jump hosts

Often, when it comes to configuring core network devices, these are isolated from the main network via a proxy or jump host. Ansible lends itself well to automating network device configuration as most of it is performed over SSH:

however, this is only helpful in a scenario where Ansible can either be installed and operated from the jump host---or, better yet, can operate via a host such as this.

Fortunately, Ansible can do exactly that. Let's assume that you have two Cumulus Networks switches in your network (these are based on a special distribution of Linux for switching hardware, which is very similar to Debian). These two switches have the [cmls01.example.com] and [cmls02.example.com] hostnames, but both can only be accessed from a host called [bastion.example.com].

The configuration to support our [bastion] host is performed in the inventory, rather than in the playbook. We begin by defining an inventory group with the switches in, in the normal manner:

```
[switches]
cmls01.example.com
cmls02.example.com
```

However, we can now start to get clever by adding some special SSH arguments into the inventory variables for this group. Add the following code to your inventory file:

```
[switches:vars]
ansible_ssh_common_args='-o ProxyCommand="ssh -W %h:%p -q bastion.example.com"'
```

This special variable content tells Ansible to add extra options when it sets up an SSH connection, including to proxy via the [bastion.example.com] host. The [-W %h:%p] options tell SSH to proxy the connection and to connect to the host specified by [%h] (this is either [cmls01.example.com] or [cmls02.example.com]) on the port specified by [%p] (usually port [22]).

Now, if we attempt to run the Ansible [ping] module against this inventory, we can see whether it works:

```
$ ansible -i switches -m ping all
cmls02.example.com | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
127.0.0.1 app02.example.com
    "ping": "pong"
}
cmls01.example.com | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
```

You will notice that we can't actually see any differences in Ansible's behavior from the command-line output. On the surface, Ansible works just as it normally does and connects successfully to the two hosts. However, behind the scenes it proxies via [bastion.example.com].

Note that this simple example assumes that you are connecting to both the [bastion] host and [switches] using the same username and SSH credentials (or in this case, keys). There are ways to provide separate credentials for both variables, but this involves more advanced usage of OpenSSH, which is beyond the scope of this course. However, this section intends to give you a starting point and demonstrate the possibility of this, and you are free to explore OpenSSH proxying by yourself.

Let's now change track and explore how it is possible to set up Ansible to prompt you for data during a playbook run.

Configuring playbook prompts

Let's reuse the two host frontend inventories we defined at the beginning of this lab. Now, let's demonstrate how to capture data from users during a playbook run with a practical example:

1. Create a simple play definition in the usual manner, as follows:

```
---
- name: A simple play to demonstrate prompting in a playbook
  hosts: frontends
```

2. Now, we'll add a special section to the play definition. We previously defined a [vars] section, but this time we will define one called [vars_prompt] (which enables you to do just that---define variables through user prompts). In this section, we will prompt for two variables---one for a user ID and one for a password. One will be echoed to the screen, while the other won't be, by setting [private: yes]:

```
vars_prompt:
  - name: loginid
    prompt: "Enter your username"
    private: no
  - name: password
    prompt: "Enter your password"
    private: yes
```

3. We'll now add a single task to our playbook to demonstrate this prompting process of setting the variables:

```
tasks:
  - name: Proceed with login
    debug:
      msg: "Logging in as {{ loginid }}..."
```

4. Now, let's run the playbook and see how it behaves:

```
$ ansible-playbook -i hosts prompt.yml
Enter your username: james
Enter your password:

PLAY [A simple play to demonstrate prompting in a playbook] ********************

TASK [Gathering Facts] ********************************************************
ok: [frt01.example.com]
ok: [frt02.example.com]

TASK [Proceed with login] *****************************************************
ok: [frt01.example.com] => {
    "msg": "Logging in as james..."
}
ok: [frt02.example.com] => {
    "msg": "Logging in as james..."
```

```
}

PLAY RECAP ***********************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

As you can see, we are prompted for both variables, yet the password is not echoed to the terminal, which is important for security reasons. We can then make use of the variables later in the playbook. Here, we just used a simple [debug] command to demonstrate that the variables have been set; however, you would instead implement an actual authentication function in place of this.

With this complete, let's proceed to the next section and look at how you can selectively run your tasks from within your plays with the use of tags.

Placing tags in the plays and tasks

In this section we will build a simple playbook with two --- each bearing a different tag to show you how tags work. We will work with the two simple host inventories that we worked with previously:

1. Create the following simple playbook to perform two tasks---one to install the [nginx] package and the other to deploy a configuration file from a template:

```
---
- name: Simple play to demonstrate use of tags
  hosts: frontends

  tasks:
    - name: Install nginx
      apt:
        name: nginx
        state: present
      tags:
        - install

    - name: Install nginx configuration from template
      template:
        src: templates/nginx.conf.j2
        dest: /etc/nginx.conf
      tags:
        - customize
```

2. Now, let's run the playbook in the usual manner, but with one difference---this time, we'll add the [--tags] switch to the command line. This switch tells Ansible to only run the tasks that have tags matching the ones that are specified. So, for example, run the following command:

```
$ ansible-playbook -i hosts tags.yml --tags install

PLAY [Simple play to demonstrate use of tags] *********************************

TASK [Gathering Facts] ********************************************************
```

```
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [Install nginx] *****************************************************
changed: [frt02.example.com]
changed: [frt01.example.com]

PLAY RECAP ***************************************************************
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Notice that the task to deploy the configuration file doesn't run. This is because it is tagged with [customize] and we did not specify this tag when running the playbook.

3. There is also a [--skip-tags] switch that does the reverse of the previous switch---it tells Ansible to skip the tags listed. So, if we run the playbook again but skip the [customize] tag, we should see an output similar to the following:

```
$ ansible-playbook -i hosts tags.yml --skip-tags customize

PLAY [Simple play to demonstrate use of tags] ********************************

TASK [Gathering Facts] ****************************************************
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [Install nginx] *****************************************************
ok: [frt02.example.com]
ok: [frt01.example.com]

PLAY RECAP ***************************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

This play run is identical because, rather than including only the [install]-tagged tasks, we skipped the tasks tagged with [customize].

Note that if you don't specify either [--tags] or [--skip-tags], then all the tasks are run, regardless of their tag.

A few notes about tags---first of all, each task can have more than one tag, so we see them specified in a YAML list format. If you use the [--tags] switch, a task will run if any of it's tags match the tag that was specified on the command line. Secondly, tags can be reused, so we could have five tasks that are all tagged [install], and all five tasks would be performed or skipped if you requested them to do so via [--tags] or [--skip-tags], respectively.

You can also specify more than one tag on the command line, running all the tasks that match any of the specified tags. Although the logic behind tags is relatively simple, it can take a little while to get used to it and the last thing you want to do is run your playbook on real hosts to check whether you understand tagging! A great way to figure this out is to add [--list-tasks] to your command, which---rather than running the playbook---lists the tasks from the playbook that would perform if you run it. Some examples are provided for you in the following code block, based on the example playbook we just created:

```
$ ansible-playbook -i hosts tags.yml --skip-tags customize --list-tasks

playbook: tags.yml

  play #1 (frontends): Simple play to demonstrate use of tags TAGS: []
    tasks:
      Install nginx TAGS: [install]

$ ansible-playbook -i hosts tags.yml --tags install,customize --list-tasks

playbook: tags.yml

  play #1 (frontends): Simple play to demonstrate use of tags TAGS: []
    tasks:
      Install nginx TAGS: [install]
      Install nginx configuration from template TAGS: [customize]

$ ansible-playbook -i hosts tags.yml --list-tasks

playbook: tags.yml

  play #1 (frontends): Simple play to demonstrate use of tags TAGS: []
    tasks:
      Install nginx TAGS: [install]
      Install nginx configuration from template TAGS: [customize]
```

As you can see, not only does [--list-tasks] show you which tasks would run, it also shows you which tags are associated with them, which helps you further understand how tagging works and ensure that you achieve the playbook flow that you wanted. Tags are an incredibly simple yet powerful way to control which parts of your playbook run and often when it comes to creating and maintaining large playbooks, it is better to be able to run only selected parts of the playbook at once. From here, we will move on to the final section of this lab, where we will look at securing your variable data at rest by encrypting it with Ansible Vault.

Securing data with Ansible Vault

Ansible Vault is a tool included with Ansible that allows you to encrypt your sensitive data at rest, while also using it in a playbook. Often, it is necessary to store login credentials or other sensitive data in a variable to allow a playbook to run unattended. However, this risks exposing your data to people who might use it with malicious intent. Fortunately, Ansible Vault secures your data at rest using AES-256 encryption, meaning your sensitive data is safe from prying eyes.

Let's proceed with a simple example that shows you how you can use Ansible Vault:

1. Start by creating a new vault to store sensitive data in; we'll call this file [secret.yml]. You can create this using the following command:

```
$ ansible-vault create secret.yml
New Vault password:
Confirm New Vault password:
```

Enter the password you have chosen for the vault when prompted and confirm it by entering it a second time (the vault that accompanies this course on GitHub is encrypted with the [secure] password).

2. When you have entered the password, you will be set to your normal editor (defined by the [EDITOR] shell variable). On my test system, this is [vi]). Within this editor, you should create a [vars] file, in the normal manner, containing your sensitive data:

```
---
secretdata: "Ansible is cool!"
```

3. Save and exit the editor (press *Esc*, then [:wq] in [vi]). You will exit to the shell. Now, if you look at the contents of your file, you will see that they are encrypted and are safe from anyone who shouldn't be able to read the file:

```
$ cat secret.yml
$ANSIBLE_VAULT;1.1;AES256
633337346237646338656332373331663336343533343738623463334643631303163653931306138
633435646539646364393632316332313237383636336461370a343236386266313331653964326334
623637376631653365396332623666363833364343663396335643562346362633664732613830
6139363035373736370a646661396464386364653935636366633663623261633538626230616630
353464653464306364633238386130373866363333333435626562396463376333353236656561323266
36646136626432633834646437346336323831383636663323730
```

4. However, the great thing about Ansible Vault is that you can use this encrypted file in a playbook as if it were a normal [variables] file (although, obviously, you have to tell Ansible your vault password). Let's create a simple playbook as follows:

```
---
- name: A play that makes use of an Ansible Vault
  hosts: frontends

  vars_files:
    - secret.yml

  tasks:
    - name: Tell me a secret
      debug:
        msg: "Your secret data is: {{ secretdata }}"
```

The [vars_files] directive is used in exactly the same way as it would be if you were using an unencrypted [variables] file. Ansible reads the headers of the [variables] files at run time and determines whether they are encrypted or not.

5. Try running the playbook without telling Ansible what the vault password is---in this instance, you should receive an error such as this:

```
$ ansible-playbook -i hosts vaultplaybook.yml
ERROR! Attempting to decrypt but no vault secrets found
```

6. Ansible correctly understands that we are trying to load a [variables] file that is encrypted with [ansible-vault], but we must manually tell it the password for it to proceed. There are a number of ways of specifying passwords for vaults (more on this in a minute), but for simplicity, try running the following command and enter your vault password when prompted:

```
$ ansible-playbook -i hosts vaultplaybook.yml --ask-vault-pass
Vault password:
```

```
PLAY [A play that makes use of an Ansible Vault] ******************************

TASK [Gathering Facts] *******************************************************
ok: [frt01.example.com]
ok: [frt02.example.com]

TASK [Tell me a secret] ******************************************************
ok: [frt01.example.com] => {
    "msg": "Your secret data is: Ansible is cool!"
}
ok: [frt02.example.com] => {
    "msg": "Your secret data is: Ansible is cool!"
}

PLAY RECAP *******************************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Success! Ansible decrypted our vault file and loaded the variables into the playbook, which we can see from the [debug] statement we created. Naturally, this defeats the purpose of using a vault, but it makes for a nice example.

This is a very simple example of what you can do with vaults. There are multiple ways that you can specify passwords; you don't have to be prompted for them on the command line---they can be provided either by a plain text file that contains the vault password or via a script that could obtain the password from a secure location at run time (think of a dynamic inventory script, only for returning a password rather than a hostname). The [ansible-vault] tool itself can also be used to edit, view, and change the passwords in a vault file, or even decrypt it and turn it back into plain text. The user guide for Ansible Vault is a great place to start for more information (https://docs.ansible.com/ansible/latest/user_guide/vault.html).

One thing to note is that you don't actually have to have a separate vault file for your sensitive data; you can actually include it inline in your playbook. For example, let's try re-encrypting our sensitive data for inclusion in an otherwise unencrypted playbook (again, use the [secure] password for the vault if you are testing the examples from the GitHub repository accompanying this course). Run the following command in your shell (it should produce an output similar to what is shown):

```
$ ansible-vault encrypt_string 'Ansible is cool!' --name secretdata
New Vault password:
Confirm New Vault password:
secretdata: !vault |
          $ANSIBLE_VAULT;1.1;AES256

34393431303339353735656236656130336664666337363732376262343837663738393465623930

33666230613063646439666665653162353131366633264310a623736643362663035373861343435

62346264313638656363323835323833633264636561366339326332356430383734653030306637

37363365336566230380a31636431383166646364353463353039333734616435663461306539643
          3331633833362666366663533346438653638303465666663313037636643564323065
Encryption successful
```

You can copy and paste the output of this command into a playbook. So, if we modify our earlier example, it would appear as follows:

```
---
- name: A play that makes use of an Ansible Vault
  hosts: frontends

  vars:
    secretdata: !vault |
          $ANSIBLE_VAULT;1.1;AES256

3439343130333935373565623665613033666466633736373237626234383766373839346562393 0

3366623061306364643966666665653162353131363633326431 0a623736643362663035373861343435

6234626431363865636333238353238336332646365613663393263332356430383734653030306637

3736333653365623038 0a31363431383316664636435346335303933373461643566634613065396434
          3331633833362663666635333346438653638303465666663313037636435643230 65

  tasks:
    - name: Tell me a secret
      debug:
        msg: "Your secret data is: {{ secretdata }}"
```

Now, when you run this playbook in exactly the same manner as we did before (specifying the vault password using a user prompt), you should see that it runs just as when we used an external encrypted [variables] file:

```
$ ansible-playbook -i hosts inlinevaultplaybook.yml --ask-vault-pass
Vault password:

PLAY [A play that makes use of an Ansible Vault] *******************************

TASK [Gathering Facts] ********************************************************
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [Tell me a secret] *******************************************************
ok: [frt01.example.com] => {
    "msg": "Your secret data is: Ansible is cool!"
}
ok: [frt02.example.com] => {
    "msg": "Your secret data is: Ansible is cool!"
}

PLAY RECAP ********************************************************************
frt01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

Ansible Vault is a powerful and versatile tool for encrypting your sensitive playbook data at rest and should enable you (with a little care) to run most of your playbooks unattended without ever leaving passwords or other sensitive

data in the clear. That concludes this section and this lab; I hope that it has been useful for you.

Summary

In this lab, you learned about running tasks asynchronously in Ansible, before looking at the various features available for running playbooks to upgrade a cluster, such as running tasks on small batches of inventory hosts, failing a play early if a certain percentage of hosts fail, delegating tasks to a specific host, and even running tasks once, regardless of your inventory (or batch) size. You also learned about the difference between running playbooks locally as opposed to on [localhost] and how to use SSH-proxying to automate tasks on an isolated network via a [bastion] host. Finally, you learned about handling sensitive data without storing it unencrypted at rest, either through prompting the user at run time or through the use of Ansible Vault. You even learned about running a subset of your playbook tasks with tagging.

In the next lab, we will explore a topic we touched on briefly in this lab in more detail---automating network device management with Ansible.

Questions

1. Which parameter allows you to configure the maximum number of hosts in a batch that will fail before a play is aborted?

A) [percentage]

B) [max_fail]

C) [max_fail_percentage]

D) [max_percentage]

E) [fail_percentage]

2. True or false -- you can use the [--connect=local] parameter to run any playbooks locally without using SSH:

A) True

B) False

3. True or false -- in order to run a playbook asynchronously, you need to use the [async] keyword:

A) True

B) False